



(Faster) Non-Maximum Suppression in Python

by Adrian Rosebrock on February 16, 2015 in [Machine Learning](#), [Tutorials](#)



I have issues — *I can't stop thinking about object detection.*

You see, last night I was watching *The Walking Dead* and instead of enjoying the zombie brutality, the forced cannibalism, or the enthralling storyline, ***all I wanted to do was build an object detection system to recognize zombies.***

Would it be very useful? Probably not.

I mean, it's quite obvious if a zombie is coming after you. The stench alone should be a ***dead*** (hey, look at that pun) giveaway, let alone the gnashing of teeth and outstretched arms. And we might as well throw in guttural "*brainnnnsss*" moans as well.

Like I said, it would be pretty obvious if a zombie was after you — you certainly wouldn't need a computer vision system to tell you that. But this is just an example of the type of stuff that runs through my head on a day-to-day basis.

To give you some context, two weeks ago I posted about using [Histogram of Oriented Gradients](#) and a [Linear Support Vector Machine](#) to build an object detection system. Sick of using the OpenCV Haar cascades and obtaining poor performance, and not to mention, extremely long training times, I took it upon myself to write my own Python object detection framework.

So far, it's gone extremely well and it's been a lot of fun to implement.

But there's *an unescapable issue* you must handle when building an object detection system — *overlapping bounding boxes*. It's going to happen, there's no way around it. In fact, it's actually a good sign that your object detector is firing properly so I wouldn't even call it an "issue" exactly.

To handle the removal overlapping bounding boxes (that refer to the same object) we can either use non-maximum suppression on the [Mean-Shift algorithm](#). While [Dalal and Triggs](#) prefer Mean-Shift, I find Mean-Shift to give sub-par results.

After receiving advice from my friend (and expert on object detection), [Dr. Tomasz Malisiewicz](#), I decided to port his non-maximum suppression MATLAB implementations over to Python.

[Last week](#) I showed you how to implement the Felzenszwalb et al. method. And this week I am going to show you the [Malisiewicz et al. method](#) which is over *100x faster*.

Note: *I intended on publishing this blog post way back in November, but due to some poor planning on my part, it's taken awhile for me to get this post out the door. Anyway, it's online now!*

So where does the speedup come from? And how do we obtain such faster suppression times?

Read on to find out.

Looking for the source code to this post?
[Jump right to the downloads section.](#)

OpenCV and Python versions:

This example will run on **Python 2.7/Python 3.4+** and **OpenCV 2.4.X/OpenCV 3.0+**.

(Faster) Non-Maximum Suppression in Python

Before we get started, if you haven't read [last week's post on non-maximum suppression](#), I would definitely start there.

Otherwise, open up a new file in your favorite editor, name it `nms.py`, and let's get started on creating a faster non-maximum suppression implementation:

Faster Non-Maximum Suppression in Python	Python
<pre>1 # import the necessary packages 2 import numpy as np 3 4 # Malisiewicz et al. 5 def non_max_suppression_fast(boxes, overlapThresh): 6 # if there are no boxes, return an empty list 7 if len(boxes) == 0: 8 return [] 9</pre>	

```

10 # if the bounding boxes integers, convert them to floats --
11 # this is important since we'll be doing a bunch of divisions
12 if boxes.dtype.kind == "i":
13     boxes = boxes.astype("float")
14
15 # initialize the list of picked indexes
16 pick = []
17
18 # grab the coordinates of the bounding boxes
19 x1 = boxes[:,0]
20 y1 = boxes[:,1]
21 x2 = boxes[:,2]
22 y2 = boxes[:,3]
23
24 # compute the area of the bounding boxes and sort the bounding
25 # boxes by the bottom-right y-coordinate of the bounding box
26 area = (x2 - x1 + 1) * (y2 - y1 + 1)
27 idxs = np.argsort(y2)
28
29 # keep looping while some indexes still remain in the indexes
30 # list
31 while len(idxs) > 0:
32     # grab the last index in the indexes list and add the
33     # index value to the list of picked indexes
34     last = len(idxs) - 1
35     i = idxs[last]
36     pick.append(i)
37
38     # find the largest (x, y) coordinates for the start of
39     # the bounding box and the smallest (x, y) coordinates
40     # for the end of the bounding box
41     xx1 = np.maximum(x1[i], x1[idxs[:last]])
42     yy1 = np.maximum(y1[i], y1[idxs[:last]])
43     xx2 = np.minimum(x2[i], x2[idxs[:last]])
44     yy2 = np.minimum(y2[i], y2[idxs[:last]])
45
46     # compute the width and height of the bounding box
47     w = np.maximum(0, xx2 - xx1 + 1)
48     h = np.maximum(0, yy2 - yy1 + 1)
49
50     # compute the ratio of overlap
51     overlap = (w * h) / area[idxs[:last]]
52
53     # delete all indexes from the index list that have
54     idxs = np.delete(idxs, np.concatenate(([last],
55         np.where(overlap > overlapThresh)[0])))
56
57 # return only the bounding boxes that were picked using the
58 # integer data type
59 return boxes[pick].astype("int")

```

Take a second to really examine this code and compare it to the Felzenszwalb et al. implementation of last week.

The code looks *almost identical*, right?

So you're probably asking yourself "Where does this 100x speed-up come from?"

And the answer is the removal of an inner `for` loop.

The implementation from last week required an extra inner `for` loop to compute the size of bounding regions and compute the ratio of overlapped area.

Instead, Dr. Malisiewicz replaced this inner `for` loop with vectorized code — *and this is how we are able to achieve substantially faster speeds when applying non-maximum suppression.*

Instead of repeating myself and going line-by-line through the code like I did last week, let's just examine the important parts.

Lines 6-22 of our faster non-maximum suppression function are essentially the same as they are from last week. We start by grabbing the (x, y) coordinates of the bounding boxes, computing their area, and sorting the indexes into the `boxes` list according to their bottom-right y-coordinate of each box.

Lines 31-55 contain our speed-up, where **Lines 41-55** are especially important.

Instead of using an inner `for` loop to loop over each of the individual boxes, we instead vectorize the code using the `np.maximum` and `np.minimum` functions — this allows us to find the maximum and minimum values across the *axis* rather than just *individual scalars*.

Note: You *have* to use the `np.maximum` and `np.minimum` functions here — they allow you to mix scalars and vectors. The `np.max` and `np.min` functions do not, and if you use them, you will find yourself with some pretty fierce bugs to find and fix. It took me quite awhile to debug this problem when I was porting the algorithm from MATLAB to Python.

Lines 47 and 48 are also vectorized — here we compute the width and height of each of the rectangles to check. Similarly, computing the `overlap` ratio on **Line 51** is also vectorized. From there, we just delete all entries from our `idx` list that are greater than our supplied overlap threshold. Typical values for the overlap threshold normally fall in the range 0.3-0.5.

The Malisiewicz et al. method is essentially identical to the Felzenszwalb et al. method, but by using vectorized code we are able to reach a [reported 100x speed-up](#) in non-maximum suppression!

Faster Non-Maximum Suppression in Action

Let's go ahead and explore a few examples. We'll start with the creepy little girl zombie that we saw at the top of this image:



Figure 1: Detecting **three** bounding boxes in the image, but non-maximum suppression suppresses two of the overlapping boxes.

It's actually really funny how well our face detector trained on real, healthy human faces generalizes to zombie faces as well. Granted, they are still "human" faces, but with all the blood and disfigurement, I wouldn't be surprised to see some strange results.


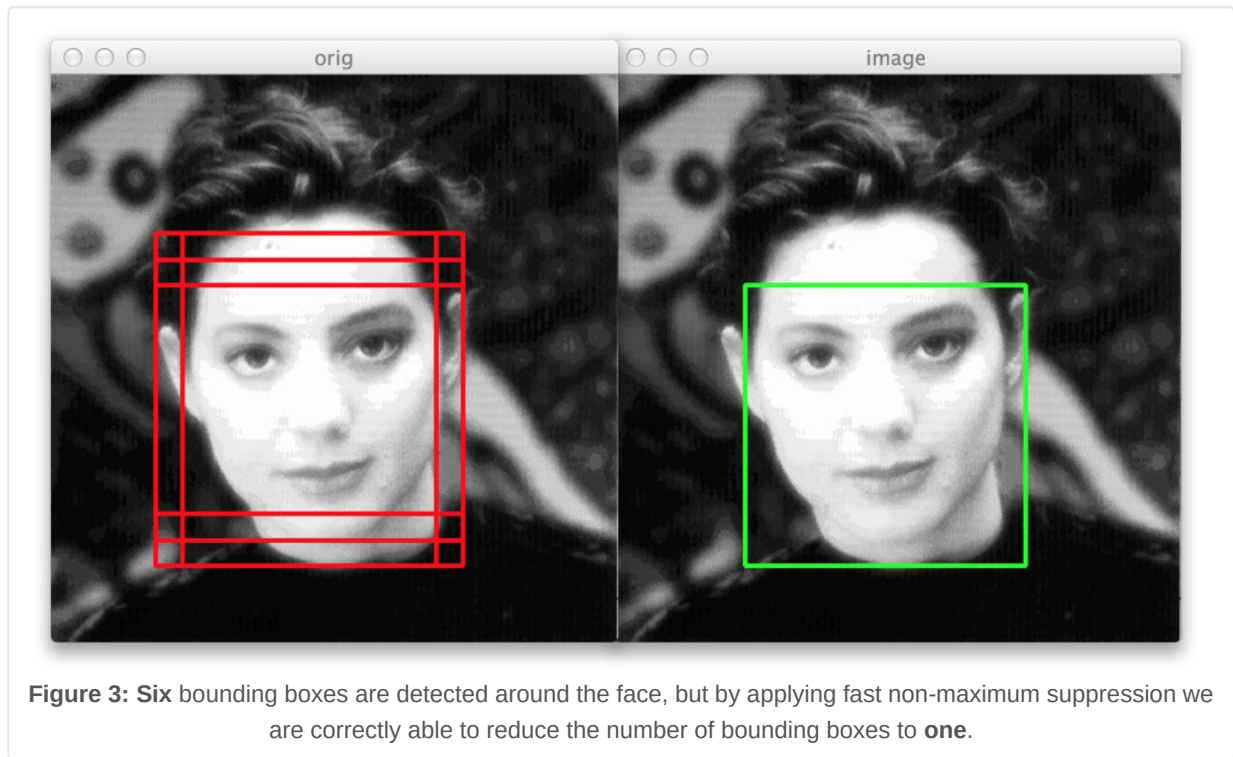


Figure 2: It looks like our face detector didn't generalize as well — the teeth of this zombie looks like a face to the detector.

Speaking of strange results, it looks like our face detector detected the mouth/teeth region of the zombie on the right. Perhaps if I had explicitly trained the HOG + Linear SVM face detector on zombie images the results would be better.



In this last example we can again see that our non-maximum suppression algorithm is working correctly — even though **six** original bounding boxes were detected by the HOG + Linear SVM detector, applying non-maximum suppression has correctly suppressed five of the boxes, leaving us with the final detection.

Summary

In this blog post we reviewed the Malisiewicz et al. method for non-maximum suppression.

This method is nearly identical to the Felzenszwalb et al. method, but by removing an inner `for` loop and using vectorized code we are able to obtain substantial speed-ups.

And if you have a second, [be sure to thank Dr. Tomasz Malisiewicz!](#)

Downloads:



If you would like to download the code and images used in this post, please enter your email address in the form below. Not only will you get a .zip of the code, I'll also send you a **FREE 17-page Resource Guide on Computer Vision, OpenCV, and Deep Learning**. Inside you'll find my hand-picked tutorials, books, courses, and libraries to help you master CV and DL! Sound good? If so, enter your email address and I'll send you the code immediately!

Email address: